

PROFESSIONAL SUMMARY

AI / Data Engineer with **11+ years** of experience, now focused on building production Generative AI solutions and agentic workflows on Databricks — LLMs, RAG pipelines, prompt design, deployment and evaluation on top of governed enterprise data. Databricks Certified Generative AI Engineer & Data Engineer Associate, backed by deep expertise in Snowflake, PySpark, Python and SQL, with a strong track record in large-scale ETL/ELT, cloud-native data engineering, performance tuning and data governance.

PROFESSIONAL EXPERIENCE

Principal Data Engineer → AI Developer · Genpact, Noida Jan 2024 — Present

- Delivering data engineering and AI solutions for a leading global financial services & investment banking client, across their enterprise data and AI initiatives.
- Built scalable pipelines ingesting enterprise (SAP) and operational sources into Snowflake via dbt-driven ELT, with PySpark transforms (JSON, Parquet), governance and quality across schemas.
- Designed scheduled & event-triggered orchestration with monitoring, retry and alerting; optimized Snowflake SQL/warehouse performance to cut compute cost.
- Transitioned into an AI Developer role — building end-to-end GenAI solutions and AI agents on Databricks using LLMs, the Databricks GenAI stack, Snowflake and PyTest.
- Developed production-grade agentic workflows — prompt design, RAG pipelines, deployment, evaluation and monitoring; led R&D across LLMs, LangChain and Ollama to prototype new agent architectures.

AWS Data Engineer · Tata Consultancy Services (TCS), Noida Jan 2023 — Jan 2024

- Delivered the Green Grass IT/OT data service — ingesting and processing industrial/operational data with a MySQL backend for downstream reporting and analytics.
- Built end-to-end big data pipelines on AWS with Spark/PySpark, Glue and Step Functions, and an ingestion layer with S3, Python, SQS, Kafka, Lambda and dbt landing curated data into Snowflake.
- Implemented event-driven orchestration with error handling, retries and monitoring for reliable production runs.

Senior Data Engineer · Biz2Credit, Noida Jun 2019 — Dec 2022

- Built and maintained data pipelines for a fintech lending platform (small-business loans) using Python (AWS Lambda), S3, Redshift and Snowflake.
- Owned reporting & analytics with Power BI on a Redshift warehouse, plus Google Analytics and complex SQL for business insights.
- Designed credit-risk models with machine learning and deep learning for underwriting and fraud prevention; delivered risk scorecards across Power BI, Tableau and Intellicus.

Senior Software Developer · ValueFirst Digital Media, Gurgaon Sep 2018 — May 2019

- Led a Scrum team of 5; built Kafka messaging pipelines (30% latency reduction) and Spark ETL (50% faster processing), introduced early ML use cases and moved workloads onto AWS.

Senior Web Developer · Aquad Software, Delhi Aug 2017 — Sep 2018

- Backend and API development in Python (Django/Flask) for web and iOS/Android projects; integrated REST APIs and third-party services (Salesforce, Marketo, Zoom).

Full Stack / Backend Developer · Samyotech Labs Pvt. Ltd., Indore Dec 2014 — Jul 2017

- Python (Django/Flask) backend developer building REST APIs and services; led a small team across database design, data integration and end-to-end delivery.

EDUCATION

B.E., Computer Science — Rajiv Gandhi Technical University (RGPV), Bhopal	2007 — 2011
GATE Qualified — Computer Science	2013 — 2015
Diploma in System Software Design (DSSD) — C-DAC, Bangalore	2012

CERTIFICATIONS

- Certified Generative AI Engineer Associate
Databricks · valid to Aug 2028
- Certified Data Engineer Associate
Databricks · valid to Jan 2028
- Claude Code 101
Anthropic · completion

TECHNICAL SKILLS

- AI / GenAI
Generative AI, LLMs, AI Agents, RAG, LangChain, Ollama, Databricks GenAI, NLP, Deep Learning, Claude, Copilot
- Warehousing & Lakehouse
Snowflake, Databricks, Delta Lake, Redshift
- Big Data & Processing
PySpark, Apache Spark, Pandas, NumPy, dbt, Kafka, EMR
- Languages
Python, SQL, PL/SQL, JavaScript, PHP, C
- Cloud
AWS S3, Lambda, Athena, EC2, IAM, Azure Blob
- ETL / Orchestration
Airflow, Azure Data Factory, AWS Glue, Step Functions, EventBridge
- Databases & Formats
PostgreSQL, MySQL, DynamoDB, Redis, Parquet, Avro
- BI & Analytics
Power BI, Tableau, Intellicus, Google Analytics
- DevOps & Tools
Data Governance, CI/CD, Docker, Git, PyTest, Agile/Scrum
- Frameworks
FastAPI, Django, Flask